

Github URL : <https://github.com/media4everyone/Data-Science-capstone-project-report>

SpaceX Falcon 9 Analysis & Prediction

SQL Data Analytics & Web Scraping Peer Assignment

Github: <https://github.com/media4everyone/Data-Science-capstone-project-report>

Author: Amr Mohamed

Dataset: SpaceX Launch Logs & Wiki Scraped Data

Domain: Aerospace Analytics & SQL

Tools: Python, SQLite, BeautifulSoup, Pandas

Project Overview & Business Value

Understanding Commercial Space Launch Economics

- › **SpaceX Milestones:** Reusable first-stage rockets drastically lower orbit insertion costs.
- › **Cost Advantage:** Falcon 9 launches cost ~\$62M vs. competitor costs upwards of \$165M.
- › **Core Analytics Objective:** Predict whether the first stage will successfully land.
- › **Business Impact:** Precise landing predictions enable accurate launch cost estimation for competitive bidding.

Data Pipeline & Environment Setup

SQL Database & Web Scraping Architecture

- **Database Stack:** SQLite database (`my_data1.db`) integrated with `ipython-sql` magic.
- **Data Cleaning:** Filtered null dates and prepared clean structured table `SPACEXTABLE`.
- **Web Scraping Stack:** `requests` + `BeautifulSoup4` parsing Wikipedia Falcon 9 launch logs.
- **Data Extraction:** Extracted HTML table headers, flight records, booster versions, and outcomes.

Task 1: Unique Launch Sites

Identifying Primary SpaceX Launch Hubs

EXECUTED SQL QUERY:

```
SELECT DISTINCT "Launch_Site" FROM SPACEXTBL;
```

QUERYRESULT OUTPUT

CCAFS LC-40, VAFB SLC-4E, KSC LC-39A, CCAFS SLC-40

[4] : **Launch_Site**

CCAFS LC-40

VAFB SLC-4E

KSC LC-39A

CCAFS SLC-40

Task 2: Cape Canaveral Launch Records

Filtering Initial CCAFS Missions

EXECUTED SQL QUERY:

```
SELECT * FROM SPACEXTBL WHERE "Launch_Site" LIKE 'CCA%' LIMIT 5;
```

[6]:

	Date	Time (UTC)	Booster_Version	Launch_Site	Payload	PAYLOAD_MASS_KG_	Orbit	Customer	Mission_Outcome	Landing_Outcome
0	2010-06-04	18:45:00	F9 v1.0 B0003	CCAFS LC-40	Dragon Spacecraft Qualification Unit	0	LEO	SpaceX	Success	Failure (parachute)
1	2010-12-08	15:43:00	F9 v1.0 B0004	CCAFS LC-40	Dragon demo flight C1, two CubeSats, barrel of...	0	LEO (ISS)	NASA (COTS) NRO	Success	Failure (parachute)
2	2012-05-22	7:44:00	F9 v1.0 B0005	CCAFS LC-40	Dragon demo flight C2	525	LEO (ISS)	NASA (COTS)	Success	No attempt
3	2012-10-08	0:35:00	F9 v1.0 B0006	CCAFS LC-40	SpaceX CRS-1	500	LEO (ISS)	NASA (CRS)	Success	No attempt
4	2013-03-01	15:10:00	F9 v1.0 B0007	CCAFS LC-40	SpaceX CRS-2	677	LEO (ISS)	NASA (CRS)	Success	No attempt

Task 3 & 4: Payload Mass Analysis

Aggregating Payload Metrics by Customer & Booster

Task 3: Total Payload for NASA (CRS)

```
SELECT SUM("PAYLOAD_MASS__KG_") FROM SPACEXTBL  
WHERE "Customer" = 'NASA (CRS)';
```

45,596 kg

Total payload mass delivered across all NASA Commercial Resupply Services missions.

Task 4: Average Payload for Booster F9 v1.1

```
SELECT AVG("PAYLOAD_MASS__KG_") FROM SPACEXTBL  
WHERE "Booster_Version" = 'F9 v1.1';
```

2,928.4 kg

Mean payload capacity carried by the Falcon 9 v1.1 variant.

Task 3

Display the total payload mass carried by boosters launched by NASA (CRS)

```
[7]: pd.read_sql('SELECT SUM("PAYLOAD_MASS__KG_") FROM SPACEXTBL WHERE "Customer" = "NASA (CRS)"', con)
```

```
[7]:
```

	SUM("PAYLOAD_MASS__KG_")
0	45596

Task 4

Display average payload mass carried by booster version F9 v1.1

```
[15]: pd.read_sql('SELECT AVG("PAYLOAD_MASS__KG_") FROM SPACEXTBL WHERE "Booster_Version" = "F9 v1.1"', con)
```

```
[15]:
```

	AVG("PAYLOAD_MASS__KG_")
0	2928.4

Task 5 & 6: Landing Milestones & Payload Filtering

Ground Pad First Success & Drone Ship Performance

Task 5: First Ground Pad Landing Success

```
SELECT MIN("Date") FROM SPACEXTBL WHERE  
"Landing_Outcome" = 'Success (ground pad)';
```

2015-12-22

Historic milestone marking SpaceX's first successful land recovery of a orbital booster.

Task 6: Drone Ship Success (4,000 - 6,000 kg)

```
SELECT "Booster_Version" FROM SPACEXTBL WHERE  
"Landing_Outcome" = 'Success (drone ship)' AND  
"PAYLOAD_MASS__KG_" BETWEEN 4000 AND 6000;
```

4 Boosters

Boosters: F9 FT B1022, F9 FT B1026, F9 FT B1021.2, F9 FT B1031.2

Task 5

List the date when the first succesful landing outcome in ground pad was acheived.

Hint: Use min function

```
16]: pd.read_sql('SELECT MIN("Date") FROM SPACEXTBL WHERE "Landing_Outcome" = "Success (ground pad)"', con)
```

```
16]:
```

	MIN("Date")
0	2015-12-22

Task 6

List the names of the boosters which have success in drone ship and have payload mass greater than 4000 but less than 6000

```
17]: pd.read_sql('SELECT "Booster_Version" FROM SPACEXTBL WHERE "Landing_Outcome" = "Success (drone ship)" AND "PAYLOAD_MASS__KG_" BETWEEN 4000 AND 6000', con)
```

```
17]:
```

	Booster_Version
0	F9 FT B1022
1	F9 FT B1026
2	F9 FT B1021.2
3	F9 FT B1031.2

Task 7: Mission Outcome Breakdown

Evaluating Overall Flight Success Rate

EXECUTED SQL QUERY:

```
SELECT "Mission_Outcome" COUNT(*) FROM SPACEXTBL GROUP BY "Mission_Outcome";
```

QUERYRESULT OUTPUT

Success: 98 missions
Success (with minor variants): 2 missions
Failure (in flight): 1 mission (CRS-7 launch failure)

Key Analytical Insight: SpaceX achieved a >98% overall launch mission success rate during the analyzed dataset period.

Task 7

```
pd.read_sql("SELECT 'Mission_Outcome', COUNT(*) FROM SPACEXTBL GROUP BY 'Mission_Outcome'", con)
```

List the total number of successful and failure mission outcomes

```
[19]: pd.read_sql("SELECT 'Mission_Outcome', COUNT(*) FROM SPACEXTBL GROUP BY 'Mission_Outcome'", con)
```

	Mission_Outcome	COUNT(*)
0	Failure (in flight)	1
1	Success	98
2	Success	1
3	Success (payload status unclear)	1

Task 8: Maximum Payload Carrying Boosters

Subquery Analysis for Heavy Payload Missions (15,600 kg)

EXECUTED SQL QUERY:

```
SELECT DISTINCT "Booster_Version" FROM SPACEXTBL WHERE "PAYLOAD_MASS_KG_" = (SELECT  
MAX("PAYLOAD_MASS_KG_") FROM SPACEXTBL);
```

QUERYRESULT OUTPUT

F9 B5 B1048.4, F9 B5 B1049.4, F9 B5 B1051.3, F9 B5 B1056.4

F9 B5 B1048.5, F9 B5 B1051.4, F9 B5 B1049.5, F9 B5 B1060.2

F9 B5 B1058.3, F9 B5 B1051.6, F9 B5 B1060.3, F9 B5 B1049.7

The screenshot shows a Jupyter Notebook interface with a cell titled "Task 8". The cell contains the following SQL query:

```
pd.read_sql("SELECT DISTINCT 'Booster_Version' FROM SPACEXTBL WHERE 'PAYLOAD_MASS_KG_' = (SELECT MAX('PAYLOAD_MASS_KG_') FROM SPACEXTBL);", con)
```

Below the query, there is a text prompt: "List all the booster versions that have carried the maximum payload mass, using a subquery with a suitable aggregate function."

The output of the query is displayed as a table with the following data:

	Booster_Version
0	F9 B5 B1048.4
1	F9 B5 B1049.4
2	F9 B5 B1051.3
3	F9 B5 B1056.4
4	F9 B5 B1048.5
5	F9 B5 B1051.4
6	F9 B5 B1049.5

Task 9: Drone Ship Landing Failures in 2015

Date & Pattern Handling in SQLite

EXECUTED SQL QUERY:

```
SELECT SUBSTR("Date", 6, 2) AS Month, "Landing_Outcome", "Booster_Version", "Launch_Site" FROM
SPACEXTBL WHERE SUBSTR("Date", 1, 4) = '2015' AND "Landing_Outcome" = 'Failure (drone ship)';
```

QUERYRESULT OUTPUT

Month 01 (Jan 2015): F9 v1.1 B1012 | CCAFS LC-40 | Failure (drone ship)

Month 04 (Apr 2015): F9 v1.1 B1015 | CCAFS LC-40 | Failure (drone ship)

Key Analytical Insight: Early 2015 experimental drone ship landings yielded key telemetry data that perfected autonomous barge landings.

Task 9

List the records which will display the month names, failure landing_outcomes in drone ship ,booster versions, launch_site for the months in year 2015.

Note: SQLite does not support monthnames. So you need to use substr(Date, 6,2) as month to get the months and substr(Date,0,5)='2015' for year.

```
pd.read_sql('SELECT SUBSTR("Date", 6, 2) AS Month, "Landing_Outcome", "Booster_Version", "Launch_Site" FROM SPACEXTBL WHERE SUBSTR("Date", 1, 4) = '2015' AND "Landing_Outcome" = 'Failure (drone ship)')
```

```
20]:
```

	Month	Landing_Outcome	Booster_Version	Launch_Site
0	01	Failure (drone ship)	F9 v1.1 B1012	CCAFS LC-40
1	04	Failure (drone ship)	F9 v1.1 B1015	CCAFS LC-40

Task 10

Task 10: Landing Outcome Ranking (2010–2017)

Historical Distribution of First Stage Recoveries

EXECUTED SQL QUERY:

```
SELECT "Landing_Outcome", COUNT(*) AS Count FROM SPACEXTBL WHERE "Date" BETWEEN '2010-06-04' AND '2017-03-20' GROUP BY "Landing_Outcome" ORDER BY Count DESC;
```

QUERYRESULT OUTPUT

1. No attempt: 10 missions
2. Success (drone ship): 5 missions
3. Failure (drone ship): 5 missions
4. Success (ground pad): 3 missions
5. Controlled (ocean): 3 missions
6. Uncontrolled (ocean): 2 missions | Failure (parachute): 2 missions

Task 10

Rank the count of landing outcomes (such as Failure (drone ship) or Success (ground pad)) between the date 2010-06-04 and 2017-03-20, in descending order.

```
[21]: pd.read_sql('SELECT "Landing_Outcome", COUNT(*) AS Count FROM SPACEXTBL WHERE "Date" BETWEEN "2010-06-04" AND "2017-03-20" GROUP BY "Landing_Outcome" ORDER BY Count DESC')
```

[21]:

	Landing_Outcome	Count
0	No attempt	10
1	Success (drone ship)	5
2	Failure (drone ship)	5
3	Success (ground pad)	3
4	Controlled (ocean)	3
5	Uncontrolled (ocean)	2
6	Failure (parachute)	2
7	Precluded (drone ship)	1

Web Scrapping Architecture: Wikipedia Data

Automated HTML Table Parsing with BeautifulSoup

- **Target Source:** Wikipedia List of Falcon 9 and Falcon Heavy launches (Snapshot 2021).
- **HTTP Requests:** Standardized user-agent headers to fetch full rendered HTML markup.
- **Parsing Helper Functions:** Handled string formatting, date extraction, and numeric regex cleaning.
- **Table Extraction:** Isolated target `wikitable plainrowheaders` starting from index 3.
- **Data Standardization:** Converted raw HTML cells into structured Pandas DataFrames.

TASK 1: Request the Falcon9 Launch Wiki page from its URL

First, let's perform an HTTP GET method to request the Falcon9 Launch HTML page, as an HTTP response.

```
[3]: import requests
from bs4 import BeautifulSoup

# Define static_url and headers
static_url = "https://en.wikipedia.org/w/index.php?title=List_of_Falcon_9_and_Falcon_Heavy_launches&oldid=1027686922"
headers = {
    'User-Agent': 'Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/115.0.0.0 Safari/537.36'
}

# use requests.get() method with the provided static_url and headers
# assign the response to a object
response = requests.get(static_url, headers=headers)
soup = BeautifulSoup(response.text, 'html.parser')
print(soup.title)
```

<title>List of Falcon 9 and Falcon Heavy launches - Wikipedia</title>

Create a BeautifulSoup object from the HTML response

```
[5]: # Use BeautifulSoup() to create a BeautifulSoup object from a response text content
soup = BeautifulSoup(response.text, 'html.parser')
```

Print the page title to verify if the BeautifulSoup object was created properly

TASK 2: Extract all column/variable names from the HTML table header

Next, we want to collect all relevant column names from the HTML table header

Let's try to find all tables on the wiki page first. If you need to refresh your memory about BeautifulSoup, please check the external reference link towards the end of this lab

```
[9]: # Use the find_all function in the BeautifulSoup object, with element type 'table'
# Assign the result to a List called 'html_tables'
response = requests.get(static_url, headers=headers)
soup = BeautifulSoup(response.text, 'html.parser')
soup.title
html_tables = soup.find_all('table')
```

Starting from the third table is our target table contains the actual launch records.

```
[10]: # Let's print the third table and check its content
first_launch_table = html_tables[2]
print(first_launch_table)

<table class="wikitable plainrowheaders collapsible" id="mwAQ" style="width: 100%;">
<tbody id="mwmg"><tr id="mmw"><th id="mwNA" scope="col">Flight No.</th>
<th id="mwNQ" scope="col">Date and<br id="mwng"/>time (<a href="https://en.wikipedia.org/wiki/Coordinated_Universal_Time" id="mwNW" r
el="mw:WikiLink" title="Coordinated Universal Time">UTC</a></th>
<th id="mwOA" scope="col"><a href="https://en.wikipedia.org/wiki/List_of_Falcon_9_first-stage_boosters" id="mwOQ" rel="mw:WikiLink" t
itle="List of Falcon 9 first-stage boosters">Version,<br id="mwog"/>Booster</a> <sup about="#mw70" class="mw-ref reference" data-mw
= {"name": "ref", "attrs": {"name": "booster", "group": "lower-alpha"}, "body": {"id": "mw-reference-text-cite_note-booster-11"}, "parts": [{"te
mplate": {"target": {"wt": "efn", "href": "/.Template:Efn"}, "params": {"name": {"wt": "booster"}, "1": {"wt": "Falcon 9 first-stage boosters are
designated with a construction serial number and an optional flight number when reused, e.g. B1021.1 and B1021.2 represent the two fl
ights of booster [[B1021]]. Launches using reused boosters are denoted with a recycled symbol ∆."}}}, {"i": 0}}}}' id="cite_ref-booster_
```

▼ TASK 3: Create a data frame by parsing the launch HTML tables



We will create an empty dictionary with keys from the extracted column names in the previous task. Later, this dictionary will be converted into a Pandas dataframe

```
[16]: launch_dict= dict.fromkeys(column_names)

# Remove an irrelevant column
del launch_dict['Date and time ( )']

# Let's initial the launch_dict with each value to be an empty list
launch_dict['Flight No.'] = []
launch_dict['Launch site'] = []
launch_dict['Payload'] = []
launch_dict['Payload mass'] = []
launch_dict['Orbit'] = []
launch_dict['Customer'] = []
launch_dict['Launch outcome'] = []
# Added some new columns
launch_dict['Version Booster']=[]
launch_dict['Booster landing']=[]
launch_dict['Date']=[]
launch_dict['Time']=[]
```

Next, we just need to fill up the `launch_dict` with launch records extracted from table rows.

```
customer = row[6].a.string
#print(customer)

# Launch outcome
# TODO: Append the launch_outcome into launch_dict with key `Launch outcome`
launch_outcome = list(row[7].strings)[0]
#print(launch_outcome)

# Booster Landing
# TODO: Append the launch_outcome into launch_dict with key `Booster Landing`
booster_landing = landing_status(row[8])
#print(booster_Landing)
```

```
F9 v1.07B0003.1
F9 v1.07B0004.1
F9 v1.07B0005.1
F9 v1.07B0006.1
F9 v1.07B0007.1
F9 v1.17B1003
F9 v1.1
F9 v1.1
F9 v1.1
F9 v1.1
F9 v1.1
F9 v1.1[
F9 v1.1[
```

Key Analytics & Technical Insights

Synthesis of SQL Queries & Scraped Telemetry

- **Booster Evolution:** Rapid transition from Falcon 9 v1.0 → v1.1 → FT → Block 5.
- **Reusability Curve:** Early parachute recovery failed; shifted to propulsive ocean/drone ship landings.
- **Payload Expansion:** Falcon 9 Block 5 standardized ~15.6t Starlink batch launches.
- **Cost Implications:** Successful first-stage landing rates directly correlate with commercial launch margin.

Next Steps for Binary Classification

-

A map of the United States with a red dot in California labeled "VARS" and "16 Oct 1945". The map shows major cities and states across the country.

Slide 14 of 15

```
data = pd.read_csv(URL1)
```

```
data.head()
```

FlightNumber	Date	BoosterVersion	PayloadMass	Orbit	LaunchSite	Outcome	Flights	GridFins	Reused	Legs	LandingPad	Block	ReusedCount	Serial	Long
--------------	------	----------------	-------------	-------	------------	---------	---------	----------	--------	------	------------	-------	-------------	--------	------

0	1	2010-06-04	Falcon 9	6104.959412	LEO	CCAFS SLC 40	None None	1	False	False	False	NaN	1.0	0	80003	-80.5
1	2	2012-05-22	Falcon 9	525.000000	LEO	CCAFS SLC 40	None None	1	False	False	False	NaN	1.0	0	80005	-80.5
2	3	2013-03-01	Falcon 9	677.000000	ISS	CCAFS SLC 40	None None	1	False	False	False	NaN	1.0	0	80007	-80.5
3	4	2013-09-29	Falcon 9	500.000000	PO	YAF8 SLC 4E	False Ocean	1	False	False	False	NaN	1.0	0	81003	-120.6
4	5	2013-12-03	Falcon 9	3170.000000	GTO	CCAFS SLC 40	None None	1	False	False	False	NaN	1.0	0	81004	-80.5

```
URL2 = 'https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DS0321EN-SkillsNetwork/datasets/dataset_part_3.csv'
X = pd.read_csv(URL2)
```


```
X.head(100)
```

FlightNumber	PayloadMass	Flights	Block	ReusedCount	Orbit_ES- L1	Orbit_GEO	Orbit_GTO	Orbit_HEO	Orbit_ISS	...	Serial_B1058	Serial_B1059	Serial_B1060
--------------	-------------	---------	-------	-------------	-----------------	-----------	-----------	-----------	-----------	-----	--------------	--------------	--------------

0	1.0	6104.959412	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
1	2.0	525.000000	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
2	3.0	677.000000	1.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0
3	4.0	500.000000	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
4	5.0	3170.000000	1.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	...	0.0	0.0	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
85	86.0	15400.000000	2.0	5.0	2.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	1.0
86	87.0	15400.000000	3.0	5.0	2.0	0.0	0.0	0.0	0.0	0.0	...	1.0	0.0	0.0
87	88.0	15400.000000	6.0	5.0	5.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0
88	89.0	15400.000000	3.0	5.0	2.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	1.0



```
[ ]: # Create a folium.PolyLine object using the coastline coordinates and launch site coordinate
# lines=folium.PolyLine(locations=coordinates, weight=1)
site_map.add_child(lines)
```

Your updated map with distance line should look like the following screenshot:




TODO: Similarly, you can draw a line between a launch site to its closest city, railway, highway, etc. You need to use `MousePosition` to find their coordinates on the map first

```
site_map.add_child(mouse_position)
site_map
```



Now zoom in to a launch site and explore its proximity to see if you can easily find any railway, highway, coastline, etc. Move your mouse to these points and mark down their coordinates (shown on the top-left) in order to the distance to the launch site.




From the color-labeled markers in marker clusters, you should be able to easily identify which launch sites have relatively high success rates.

[]: # TASK 3: Calculate the distances between a launch site to its proximities

```
# create a marker object with its coordinate
# and customize the Marker's icon property to indicate if this launch was succeeded or failed,
# e.g., icon=folium.Icon(color='white', icon_color=row['marker_color'])
for index, record in spacex_df.iterrows():
    # TODO: Create and add a Marker cluster to the site map
    # marker = folium.Marker(...)
    marker_cluster.add_child(marker)

site_map
```


Your updated map may look like the following screenshots:



```
folium.map.Marker(coordinate, icon=DivIcon(icon_size=(20,20),icon_anchor=(0,0), html='<div style="font-size: 12; color:#d35400;">
<b>%s</b></div>' % 'label', ))
```

```
[ ]: # Initial the map
site_map = folium.Map(location=nasa_coordinate, zoom_start=5)
# For each Launch site, add a Circle object based on its coordinate (Lat, Long) values. In addition, add Launch site name as a popup label
```


The generated map with marked launch sites should look similar to the following:




Now, you can explore the map by zoom-in/out the marked areas , and try to answer the following questions:

1. What are the launch sites marked on the map?

[1]:



[5]:




The following dataset with the name `spacex_launch_geo.csv` is an augmented dataset with latitude and longitude added for each site.

TASK 4

Create a logistic regression object then create a GridSearchCV object `logreg_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters`.

```
[28]: parameters = {'C':[0.01,0.1,1],
                  'penalty':['l2'],
                  'solver':['lbfgs']}
lr = LogisticRegression()
logreg_cv = GridSearchCV(lr, parameters, cv=10)
logreg_cv.fit(X_train, Y_train)
```

```
[28]:
```



```
[29]: parameters={'C':[0.01,0.1,1], 'penalty':['l2'], 'solver':['lbfgs']}# (1 Lasso (2 ridge
lr=LogisticRegression()
```

We output the `GridSearchCV` object for logistic regression. We display the best parameters using the data attribute `best_params_` and the accuracy on the validation data using the data attribute `best_score_`.

```
[30]: print("tuned hpyerparameters :(best parameters) ",logreg_cv.best_params_)
      print("accuracy :",logreg_cv.best_score_)

tuned hpyerparameters :(best parameters) {'C': 0.01, 'penalty': 'l2', 'solver': 'lbfgs'}
accuracy : 0.8464285714285713
```

TASK 5

Calculate the accuracy on the test data using the method `score`:

```
[31]: logreg_score = logreg_cv.score(X_test, Y_test)
      print("Logistic Regression Test Accuracy:", logreg_score)
```

TASK 6

Create a support vector machine object then create a `GridSearchCV` object `svm_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters`.


```
[33]: parameters = {'kernel':('linear', 'rbf', 'poly', 'rbf', 'sigmoid'),
                  'C': np.logspace(-3, 3, 5),
                  'gamma':np.logspace(-3, 3, 5)}

svm = SVC()
```

```
[34]: parameters = {'kernel':('linear', 'rbf', 'poly', 'sigmoid'),
                  'C': np.logspace(-3, 3, 5),
                  'gamma': np.logspace(-3, 3, 5)}

svm = SVC()
svm_cv = GridSearchCV(svm, parameters, cv=10)
svm_cv.fit(X_train, Y_train)
```

```
[34]:
```



```
[35]: print("tuned hyperparameters :(best parameters) ",svm_cv.best_params_)
print("accuracy :",svm_cv.best_score_)

tuned hyperparameters :(best parameters) {'C': np.float64(1.0), 'gamma': np.float64(0.03162277660168379), 'kernel': 'sigmoid'}
accuracy : 0.8482142857142856
```

TASK 7

Calculate the accuracy on the test data using the method `score`:

```
[36]: svm_score = svm_cv.score(X_test, Y_test)
print("SVM Test Accuracy:", svm_score)

SVM Test Accuracy: 0.8333333333333334
```

We can plot the confusion matrix

Conclusion & References

IBM Data Science Capstone & Project Summary

- **Summary:** Successfully structured SpaceX launch dataset using SQLite & Web Scraping.
- **Authors & Credits:** Lakshmi Holla, Rav Ahuja, Anita Verma (IBM Corporation).
- **References:** IBM Developer Skills Network SQL Magic & Sub-queries Labs.
- **Status:** All 10 SQL tasks completed & verified against DB requirements.